

021-02-NLP-Embedding-solution

April 18, 2024

1 021-02 - NLP Embedding - Solution Notebook

- Written by Alexandre Gazagnes
- Last update: 2024-02-01

1.1 About

Context :

Let's get the party started !

Data :

You can find the dataset [here](#).

1.2 Preliminaries

1.2.1 System

These commands will display the system information:

Uncomment theses lines if needed.

```
[ ]: # pwd
```

```
[ ]: # cd ..
```

```
[ ]: # ls
```

Install various Librairies :

```
[ ]: # !pip install -r requirements.txt >> pip.log  
# !pip freeze >> pip.freeze
```

1.2.2 Import

```
[ ]: import os, sys, warnings, secrets, datetime  
import pickle  
from IPython.display import display  
  
# from joblib import dump, load
```

```
[ ]: import pandas as pd
import numpy as np
```

```
[ ]: # import matplotlib.pyplot as plt
# import seaborn as sns
import plotly.express as px
```

```
[ ]: from sklearn.base import *
from sklearn.preprocessing import *
from sklearn.impute import *
from sklearn.model_selection import *
from sklearn.decomposition import *
from sklearn.ensemble import *
from sklearn.model_selection import *
from sklearn.pipeline import *
from sklearn.feature_extraction import *
from sklearn.dummy import *
from sklearn.feature_extraction.text import *

# from lightgbm import *
# from xgboost import *

from sklearn.linear_model import *
from sklearn.ensemble import *
from sklearn.neighbors import *
```

```
[ ]: # import nltk

# import wordcloud

# from nltk.corpus import stopwords
# from nltk.corpus import words
# from nltk.tokenize import wordpunct_tokenize

# import string

import spacy

# from spacy.lang.en.stop_words import STOP_WORDS
```

```
[ ]: import gensim

from gensim.models import KeyedVectors
from gensim.downloader import load

from gensim.models.doc2vec import Doc2Vec, TaggedDocument
from gensim.parsing.preprocessing import preprocess_string
```

```
[ ]: # import transformers

[ ]: from openai import OpenAI

[ ]: import requests # HTTP client
```

1.2.3 Graphs and Settings

```
[ ]: # sns.set()

[ ]: warnings.filterwarnings("ignore")
# warnings.filterwarnings(action="once")
```

If needed we can use a TEST_MODE to run the notebook to have a very fast execution :

```
[ ]: TEST_MODE = True

[ ]: CV = 10 # number of folds for the cross val
N_JOBS = 7 # number of cpu to use for computations
FRAC = 1.0 # we keep 100% of the dataframe
DISPLAY = True # display complex viz
TEST_SIZE = 0.25 # Train vs Test %

if TEST_MODE:
    CV = 3
    N_JOBS = -1
    FRAC = 0.3
    DISPLAY = False
    TEST_SIZE = 0.5
```

1.2.4 Thrid Parties Tools

We need some Third parties :

```
[ ]: # nltk.download("punkt")
# nltk.download("stopwords")
# nltk.download("words")
```

Some string assets :

```
[ ]: # stop_words = set(stopwords.words("english"))
# punctuation = set(string.punctuation)
# word_dict = words.words()
```

We need to download spacy :

```
[ ]: # !python -m spacy download en_core_web_sm
# !python -m spacy download en_core_web_md
# !python -m spacy download en_core_web_lg
```

Word2vect :

```
[ ]: w2c = load("word2vec-google-news-300")
```

And to load spacy model :

```
[ ]: # nlp = spacy.load("en_core_web_sm")  
  
nlp = spacy.load("en_core_web_lg")
```

If you have an error please run the download command for spacy :

```
[ ]: # !python -m spacy download en_core_web_lg
```

1.2.5 Data

url of the dataset :

```
[ ]: url = "https://gist.githubusercontent.com/AlexandreGazagnes/  
↪cabe445634a092d308d17a883a305a75/raw/  
↪d2014e8a34bba3c1be3ec8936bb338fb42888f24/nlp.csv"
```

Download the dataset :

```
[ ]: df = pd.read_csv(url)  
df.head(5)
```

Keep a copy of the df :

```
[ ]: DF = df.copy()
```

```
[ ]: if TEST_MODE:  
    df = df.sample(frac=FRAC)
```

1.3 King - Men + Woman

1.3.1 With Spacy

Tokenize 'King' :

```
[ ]: king = nlp("king")  
king
```

```
[ ]: type(king)
```

Extract the vector :

```
[ ]: king_v = king.vector  
king_v
```

Length ?

```
[ ]: len(king.vector)
```

Same for Man :

```
[ ]: man = nlp("man")
man_v = man.vector
man_v
```

```
[ ]: len(man_v)
```

Same for woman :

```
[ ]: woman = nlp("woman")
woman_v = woman.vector
woman_v
```

Fancy calculation !

```
[ ]: res = king_v - man_v + woman_v
res
```

Length ?

```
[ ]: len(res)
```

Reshape new vector :

```
[ ]: res = res.reshape(1, -1)
res
```

Compute Similarity :

```
[ ]: vectors = nlp.vocab.vectors.most_similar(res, n=20)
vectors
```

v1 is :

```
[ ]: v1 = vectors[0][0][0]
v1
```

vect is :

```
[ ]: vect = nlp.vocab[v1]
vect
```

text is :

```
[ ]: vect.text
```

```
[ ]: v2 = vectors[0][0][1]
vect = nlp.vocab[v2]
vect.text
```

```
[ ]: v3 = vectors[0][0][2]
vect = nlp.vocab[v3]
vect.text
```

Whoooo not so good !

Lets do the same with a “huge” model :

```
[ ]: # !python -m spacy download en_core_web_sm
# !python -m spacy download en_core_web_md
!python -m spacy download en_core_web_lg
# !python -m spacy download en_core_web_trf
```

```
[ ]: nlp = spacy.load("en_core_web_lg")
```

Good ? ...

Just re-run previous cells with this code.

What are your conclusions ?

Let's try another last trick :

```
[ ]: doc = "He is one of the most famous kings: Richard III was the last king of_
↪England to die in battle"
doc = nlp(doc)
king = doc[-7]
king
```

```
[ ]: king_v = king.vector
```

```
[ ]: doc = "Fifteen months after the death of King George VI, his daughter Elizabeth_
↪is crowned Queen of England"
doc = nlp(doc)
queen = doc[-3]
queen
```

```
[ ]: queen_v = queen.vector
```

```
[ ]: doc = "a female, it's a woman, or a lady, a human of female sex."
doc = nlp(doc)
woman = doc[6]
woman
```

```
[ ]: woman_v = woman.vector
```

```
[ ]: king_v = king.vector
```

```
[ ]: doc = "a boy, a guy, or a man, it's a human being of male sex."
doc = nlp(doc)
```

```
man = doc[8]
man
```

```
[ ]: man_v = man.vector
```

```
[ ]: out = nlp.vocab.vectors.most_similar(queen_v.reshape(1, -1), n=20)
out
```

```
[ ]: for t_id in out[0][0]:
    print(nlp.vocab[t_id].text)
```

1.3.2 With Doc2Vect

Let's do the same with Pretrained Doc2Vect :

```
[ ]: result = w2c.most_similar(positive=["woman", "king"], negative=["man"], topn=10)
result
```

1.4 Using Gensim

1.4.1 Prepare Data

Create y vector :

```
[ ]: y = df.cat_1
y
```

Create X :

```
[ ]: X = df.description
X
```

Cross validation :

```
[ ]: def cv():
    return StratifiedShuffleSplit(n_splits=CV, test_size=TEST_SIZE)

cv()
```

1.4.2 By Hand

Our documents :

```
[ ]: documents = df.description
documents[:10]
```

Init spacy :

```
[ ]: nlp = spacy.load("en_core_web_lg")
```

If you have an error please download en_core_web_lg model for spacy

Preprocess (clean) the corpus :

```
[ ]: tokenized_docs = [
    [
        token.lemma_
        for token in nlp(doc.lower())
        if not token.is_stop and not token.is_punct
    ]
    for doc in documents
]
tokenized_docs[:10]
```

```
[ ]: tokenized_docs[0]
```

Key concept here is a tagged document => Token + id

```
[ ]: tagged_docs = [
    TaggedDocument(words=doc, tags=[i]) for i, doc in enumerate(tokenized_docs)
]
tagged_docs[:10]
```

1.5 Train the Doc2Vec model

1.5.1 Building Gensim Models

sm :

```
[ ]: # 5s
model_sm = Doc2Vec(
    tagged_docs,
    vector_size=50, # size of output vect
    window=2, # nb words before and after a target word
    min_count=1, # minimum frequency count of words. ,
    workers=4, # number of cpu
    epochs=100, # number of iterations (passes over the entire dataset)
)

model_sm
```

md :

```
[ ]: # 10s
model_md = Doc2Vec(
    tagged_docs,
    vector_size=100,
    window=4,
    min_count=1,
    workers=4,
```



```
        epochs=200,  
    )  
    model_md
```

lg :

```
[ ]: # 30s  
model_lg = Doc2Vec(  
    tagged_docs,  
    vector_size=500,  
    window=10,  
    min_count=1,  
    workers=4,  
    epochs=500,  
)  
model_lg
```

xl :

```
[ ]: # # 15m => 1h  
# model_xxl = Doc2Vec(  
#     tagged_docs,  
#     vector_size=1_000,  
#     window=10,  
#     min_count=3,  
#     workers=6,  
#     epochs=2_000,  
# )  
# model_xxl
```

```
[ ]: # 1m => 3m  
model_xl = Doc2Vec(  
    tagged_docs,  
    vector_size=1_000,  
    window=10,  
    min_count=1,  
    workers=4,  
    epochs=1_000,  
)  
model_xl
```

Get the vectors :

```
[ ]: # 2 => 5 mins with xl  
  
doc_vectors = []  
for doc in tokenized_docs:  
    vector = model_xl.infer_vector(doc)  
    doc_vectors.append(vector)
```

```
doc_vectors
```

A more pythonic implementation

```
[ ]: doc_vectors = [model_xl.infer_vector(doc) for doc in tokenized_docs]
doc_vectors
```

Data Type :

```
[ ]: type(doc_vectors)
```

```
[ ]: type(doc_vectors[0])
```

Length ? :

```
[ ]: len(doc_vectors)
```

```
[ ]: df.shape
```

```
[ ]: len(doc_vectors[0])
```

```
[ ]: len(df)
```

Rebuild a 'special' X :

```
[ ]: X = pd.DataFrame(doc_vectors)
X
```

```
[ ]: X.to_csv("df_from_doc2vect_model_xl.csv", index=False)
```

Shape :

```
[ ]: X.shape
```

Grid :

```
[ ]: grid = GridSearchCV(
    RandomForestClassifier(),
    {},
    cv=CV,
    n_jobs=N_JOBS,
    return_train_score=True,
    verbose=1,
)
grid
```

```
[ ]: grid.fit(X, y)
```

```
[ ]: display(grid.best_estimator_)
```

Resultize :

```
[ ]: def resultize(grid, head=20, export=False, token=None):

    res = grid.cv_results_
    res = pd.DataFrame(res)

    cols = [i for i in res.columns if "split" not in i]
    res = res.loc[:, cols]

    res = res.drop(columns=["mean_score_time", "std_score_time"])

    res = res.round(2).sort_values("mean_test_score", ascending=False)

    if export:
        _res = res.copy().head(head)
        _res["token"] = token
        _res = _res.astype(str)
        now = str(datetime.datetime.now())[19].replace(" ", "_")
        _res.to_csv(f"results_{token}_{now}.csv", index=False)

    return res.head(head)
```

```
[ ]: resultize(grid)
```

1.5.2 Using a pipeline

```
[ ]: pipeline = Pipeline(
    [
        ("preprocessor", "passthrough"),
        ("scaler", "passthrough"),
        ("reductor", "passthrough"),
        ("estimator", LogisticRegression()),
    ]
)

pipeline
```

What is “passthrough” :

```
[ ]: pst = "passthrough"

pst
```

Param grid :

```
[ ]: param_grid = {
    "scaler": [
```

```

        "passthrough",
        StandardScaler(),
        QuantileTransformer(n_quantiles=100),
        Normalizer(),
    ],
    "reductor": [PCA()],
    "reductor__n_components": [0.7, 0.85, 0.9, 0.95, 0.99],
    "estimator": [RandomForestClassifier(), LogisticRegression()],
}
param_grid

```

New grid :

```

[ ]: grid = GridSearchCV(
    pipeline,
    param_grid=param_grid,
    cv=CV,
    n_jobs=N_JOBS,
    return_train_score=True,
    verbose=1,
)
grid

```

```

[ ]: grid.fit(X, y)

```

Results

```

[ ]: display(grid.best_estimator_)

```

```

[ ]: resultsize(grid)

```

1.5.3 Using a custom transformer

Our Transformer :

```

[ ]: class Doc2VecTransformer(BaseEstimator, TransformerMixin):
    """ """

    def __init__(
        self,
        model=None,
        vector_size=500,
        window=5,
        min_count=5,
        epochs=100,
    ):

        self.vector_size = vector_size
        self.window = window

```

```

self.min_count = min_count
self.epochs = epochs
self.model = model

def fit(self, X, y=None):

    if not isinstance(X, list):
        _X = X.values.tolist()
    else:
        _X = X

    if self.model:
        return self

    tagged_docs = [
        TaggedDocument(words=preprocess_string(doc), tags=[i])
        for i, doc in enumerate(_X)
    ]
    model = Doc2Vec(
        vector_size=self.vector_size, min_count=self.min_count, epochs=self.
↪epochs
    )
    model.build_vocab(tagged_docs)
    model.train(tagged_docs, total_examples=model.corpus_count,
↪epochs=model.epochs)
    self.model = model

    return self

def transform(self, X, y=None):

    if not isinstance(X, list):
        _X = X.values.tolist()
    else:
        _X = X

    vectors = [self.model.infer_vector(preprocess_string(i)) for i in X]
    return vectors

```

Original df :

```
[ ]: df
```

Init d2f :

```
[ ]: d2v = Doc2VecTransformer()
d2v
```

Fit :

```
[ ]: d2v.fit(df.description)
```

Transform :

```
[ ]: text = ["my new watch is a very funny flic flac digital chronometer"]
vector_list = d2v.transform(text)
vector_list
```

To be sure :

```
[ ]: vector_list[0]
```

With pretrained model :

```
[ ]: text = ["my new watch is a very funny flic flac digital chronometer"]
d2v = Doc2VecTransformer(model=model_xl)
d2v.transform(text)
```

New param grid :

```
[ ]: param_grid = {"preprocessor": [Doc2VecTransformer()]}
```

```
[ ]: pipeline
```

New Grid :

```
[ ]: grid = GridSearchCV(
    pipeline,
    param_grid,
    cv=CV,
    n_jobs=N_JOBS,
    return_train_score=True,
    verbose=2,
)
grid
```

Fit :

```
[ ]: grid.fit(df.description.values, y)
```

```
[ ]: grid.best_estimator_
```

Results :

```
[ ]: resultize(grid)
```

Just as a remainder :

```
[ ]: # # 30s
# model_lg = Doc2Vec(
#     tagged_docs,
```

```
#     vector_size=500,
#     window=10,
#     min_count=1,
#     workers=4,
#     epochs=500,
# )
# model_lg
```

Testing various transformers params :

```
[ ]: param_grid = {
    "preprocessor": [Doc2VecTransformer()],
    "preprocessor__vector_size": [500], # 100, 200, # 1000
    "preprocessor__window": [5], # 5, 10, 5, 10, # 20, 25, 30
    "preprocessor__epochs": [500],
    "preprocessor__min_count": [1, 3], # 100, 300, # 1000
    "estimator": [LogisticRegression(), RandomForestClassifier()],
    # preprocessor__model = [model_sm, model_md, model_lg, model_xl]
}
```

```
[ ]: # param_grid = {
#     "preprocessor": [Doc2VecTransformer()],
#     "preprocessor__vector_size": [500, 1000], # 100, 200, # 1000
#     "preprocessor__window": [5, 10, 15], # 5, 10, 5, 10, # 20, 25, 30
#     "preprocessor__epochs": [500, 1000],
#     "preprocessor__min_count": [1, 3, 5], # 100, 300, # 1000
#     # preprocessor__model = [model_sm, model_md, model_lg, model_xl]
# }
```

```
[ ]: # param_grid = {
#     "preprocessor": [Doc2VecTransformer()],
#     "preprocessor__vector_size": [500], # 100, 200, # 1000
#     "preprocessor__window": [5, 10], # 5, 10, 5, 10, # 20, 25, 30
#     "preprocessor__epochs": [500],
#     "preprocessor__min_count": [1, 3, 5], # 100, 300, # 1000
#     # preprocessor__model = [model_sm, model_md, model_lg, model_xl]
# }
```

```
[ ]: # param_grid = {
#     "preprocessor": [Doc2VecTransformer()],
#     "preprocessor__vector_size": [500, 1000], # 100, 200, # 1000
#     "preprocessor__window": [10, 15], # 5, 10, 5, 10, # 20, 25, 30
#     "preprocessor__epochs": [500, 1000],
#     "preprocessor__min_count": [1, 3, 5], # 100, 300, # 1000
#     # preprocessor__model = [model_sm, model_md, model_lg, model_xl]
# }
```

Grid :

```
[ ]: grid = GridSearchCV(
    pipeline,
    param_grid,
    cv=CV,
    n_jobs=N_JOBS,
    return_train_score=True,
    verbose=2,
)
```

Fit :

```
[ ]: grid.fit(df.description, y)
```

```
[ ]: grid.best_estimator_
```

Results :

```
[ ]: res = resultsize(grid, head=30)
res
```

1.5.4 Post mortem analysis

Our problem is :

```
[ ]: px.scatter_3d(
    res,
    x="param_preprocessor__vector_size",
    y="mean_test_score",
    z="param_preprocessor__window",
)
```

With box plots :

```
[ ]: px.box(
    res,
    x="param_preprocessor__vector_size",
    y="mean_test_score",
    # color="param_preprocessor__window",
)
```

```
[ ]: px.scatter(
    res.loc[res.param_preprocessor__window == 15],
    x="mean_score_time",
    y="mean_test_score",
    # color="param_preprocessor__window",
)
```


1.5.5 Exporting our model

```
[ ]: best_pipeline = grid.best_estimator_  
best_pipeline
```

Unique id for our model :

```
[ ]: token = secrets.token_hex(4)  
token
```

Date :

```
[ ]: now = str(datetime.datetime.now())[:19].replace(" ", "_")  
now
```

Filename :

```
[ ]: fn = f"gensim_model_{token}_{now}.pk"  
fn
```

Saving with pickle :

```
[ ]: with open(fn, "wb") as f:  
    pickle.dump(best_pipeline, f)
```

Size of our model :

```
[ ]: sys.getsizeof(best_pipeline)
```

```
[ ]: sys.getsizeof(grid)
```

With a function :

```
[ ]: def save(model, base_fn, token=None, score=None):  
    """Saving our model"""  
  
    if not token:  
        token = secrets.token_hex(4)  
  
    now = str(datetime.datetime.now())[:19].replace(" ", "_")  
  
    fn = f"{base_fn}_{token}_{now}"  
    if score:  
        fn += "_" + str(round(score, 2))  
  
    with open(fn + ".pk", "wb") as f:  
        pickle.dump(model, f)  
  
    return fn, sys.getsizeof(model)
```

Save our model :

```
[ ]: out = save(best_pipeline, "gensim_model")
out
```

Save our Transformer :

```
[ ]: out = save(Doc2VecTransformer, "d2v_trasnformer")
out
```

Just to be sure :

```
[ ]: best_pipeline.predict(["i have a beautiful analogic watch rolex"])
```

```
[ ]: grid.predict(["i have a beautiful analogic watch rolex"])
```

1.5.6 Load

```
[ ]: def load(fn):

    if not os.path.isfile(fn):
        raise AttributeError(f"The File {fn} do not exists!")

    with open(fn, "rb") as f:
        model = pickle.load(f)

    return model, sys.getsizeof(model)
```

List of .pk files :

```
[ ]: fn_list = os.listdir()
fn_list
```

```
[ ]: [i for i in fn_list if ".pk" in i]
```

Loading a model :

```
[ ]: fn = "-- THE FILE NAME OF YOUR MODEL --"

loaded_pipeline, _ = load(fn)
loaded_pipeline
```

Testing our loaded model :

```
[ ]: loaded_pipeline.predict(["i have a beautiful analogic watch rolex"])
```

1.6 Using OpenAI GPT Emedding

Init your client :

```
[ ]: client = OpenAI()
client
```

Doc :

```
[ ]: doc = df.description.iloc[0]
doc
```

Just a try :

```
[ ]: response = client.embeddings.create(
    input=doc,
    model="text-embedding-3-small",
)
```

What is response :

```
[ ]: response
```

The vector :

```
[ ]: vector = response.data[0].embedding
vector
```

Size?

```
[ ]: len(vector)
```

List of Vectors ?

```
[ ]: # li = []

# for doc in df.description.values :
#     # print(i)

#     response = client.embeddings.create(
#         input=doc,
#         model="text-embedding-3-small",
#     )
#     vector = response.data[0].embedding
#     li.append(vector)

# li = pd.DataFrame(li)
# li.to_csv("df_from_gpt.csv", index=False)
# li
```

Or load this file :

```
[ ]: li = pd.read_csv("df_from_gpt.csv")
li
```

With a custom transformer :

```
[ ]: class OpenAIVecTransformer(BaseEstimator, TransformerMixin):

    def __init__(self, model="text-embedding-3-small"):

        self.model = model
        self.client = OpenAI()

    def fit(self, X, y=None):

        return self

    def transform(self, X, y=None):

        if not isinstance(X, list):
            _X = X.values.tolist()
        else:
            _X = X

        get_vect = (
            lambda i: self.client.embeddings.create(
                input=i,
                model=self.model,
            )
            .data[0]
            .embedding
        )
        X_ = [get_vect(i) for i in X]

        return X_
```

Let's build a very basic Pipeline / grid search :

```
[ ]: pipeline = Pipeline(
    [
        ("preprocessor", "passthrough"),
        ("scaler", "passthrough"),
        ("reductor", "passthrough"),
        ("estimator", RandomForestClassifier()),
    ]
)

pipeline
```

What is "passthrough" :

```
[ ]: pst = "passthrough"
pst
```

Param grid :

```
[ ]: param_grid = {
    "scaler": [
        "passthrough",
        StandardScaler(),
        QuantileTransformer(n_quantiles=100),
        Normalizer(),
    ],
    # "reductor": [PCA()],
    # "reductor__n_components": [0.7, 0.85, 0.9, 0.95, 0.99],
    "estimator": [RandomForestClassifier()], # LogisticRegression()
}
param_grid
```

New grid :

```
[ ]: grid = GridSearchCV(
    pipeline,
    param_grid=param_grid,
    cv=CV,
    n_jobs=N_JOBS,
    return_train_score=True,
    verbose=1,
)

grid.fit(li, y)
```

Results

```
[ ]: display(grid.best_estimator_)
```

```
[ ]: resultize(grid)
```

1.7 Using our own API

```
[ ]: url = "https://centrale-casa-api.onrender.com"

route = "/predict/"

data = "watches"
response = requests.get(url + route + data)

response.json()
```